



BFS

observation supplies

frontier_size: 1 → 1

sail(l1,l0): new → visited

sail(l1,l2): new → new

policy emits

1 action: sail(l1,l0) · source: head

retire: true · position: 0

both: visit true · evaluate false

2 action: sail(l1,l2) · source: s0

retire: false · position: 1

Queue fields derived from frontier size and step's place in the expansion.

BFWS

observation supplies

sail [2,1] · dup: false

novelty: 1 · goals left: 1

eval.frontier: {true, 0}

step 2 [2,1]: dup:true · eval:null

policy emits

1 action: sail (l1,l0) · source: \$

both: visit/evaluate true

frontier_intent: {true, 0}

2 action: sail (l1,l2) · source: \$

frontier_intent: {false, 1}

Position supplied in candidate eval; copied, not derived by the policy.

Additive · greedy best-first

observation supplies

sail(l1,l0) g:1 h:4 p:4 → s1

sail(l1,l2) g:1 h:3 p:3 → s2

policy emits

1 action: sail(l1,l0)

source_state_id: s0

2 action: sail(l1,l2)

source_state_id: s0

Runtime computes g, h and priority.

Runtime keeps the priority heap; weighted A*: priority g + 3h.